

CONTINUOUS PROBABILITY, MARKOV CHAINS, CONDITIONAL EXPECTATION

11

COMPUTER SCIENCE MENTORS 70

April 17 - April 21, 2017

1 Conditional Expectation

1.1 Introduction

The **conditional expectation** of Y given X is defined by

$$E[Y|X = x] = \sum_y y \cdot P[Y = y|X = x] = \sum_y y \cdot \frac{P[X = x, Y = y]}{P[X = x]}$$

Properties of Conditional Expectation

$$E(a|Y) = a$$

$$E(aX + bZ|Y) = a \cdot E(X|Y) + b \cdot E(Z|Y)$$

$$E(X|Y) \geq 0 \text{ if } X \geq 0$$

$$E(X|Y) = E(X) \text{ if } X, Y \text{ independent}$$

$$E(E(X|Y)) = E(X)$$

1.2 Questions

1. Prove $E(h(X) \cdot Y|X) = h(X) \cdot E(Y|X)$

2. Prove $E(E(Y|X)) = E(Y)$

3. Consider the random variables Y and X with the following probabilities

This table gives the probability distribution for $P[X \cap Y]$

		X		
		0	1	2
Y	0	0	.1	.2
	1	.1	.2	.1
	2	.2	.1	0

Find:

(a) $E(Y|X = 0)$

(b) $E(Y|X = 1)$

(c) $E(Y|X = 2)$

(d) $E(Y)$

2 Markov Chains

P is a **transition probability matrix** if:

1. All of the entries are non-negative.
2. The sum of entries in each row is 1.

A **Markov chain** is defined by four things: $(\mathcal{X}, \pi_0, P, \{X_n\}_{n=0}^\infty)$

\mathcal{X} Set of states

π_0 Initial probability distribution

P Transition probability matrix

$\{X_n\}_{n=0}^\infty$ Sequence of random variables where:

$$P[X_0 = i] = \pi_0(i), i \in \mathcal{X}$$

$$P[X_{n+1} = j | X_n = i, X_{n-1}, \dots, X_0] = P(i, j), \forall n \geq 0, \forall i, j \in \mathcal{X}$$

A Markov chain is **irreducible** if we can go from any state to any other state, possibly in multiple steps.

Define value $d(i)$ for each state i as:

$$d(i) := \text{g.c.d}\{n > 0 | P^n(i, i) = P[X_n = i | X_0 = i] > 0\}, i \in \mathcal{X}$$

If $d(i) = 1$, then the Markov chain is **aperiodic**. If $d(i) \neq 1$, then the Markov chain is periodic and its **period** is $d(i)$.


A distribution π is **invariant** for the transition probability P if it satisfies the following **balance equations**

$$\pi \cdot P = \pi.$$

Theorem 24.3: A finite irreducible Markov chain has a unique invariant distribution.

Theorem 24.4: All irreducible and aperiodic Markov chains converge to the unique invariant distribution. If a Markov chain is finite and reducible, the amount of time spent in each state approaches the invariant distribution as n grows large

Equations that model what will happen at the next step are called **first step equations**



Denote $\beta(i, j)$ as the expected amount of time it would take to move from i to j . $\beta(0, 1) = 1 + \frac{1}{2} \cdot \beta(0, 1)$
 $\beta(1, 1) = 0$

2.1 Questions

1. Life of Alex

Alex is enjoying college life. She spends a day either studying, partying, or looking for housing for the next year. If she is studying, the chances of her studying the next day are 30%, the chances of her partying the next day are 50%, and the chances of her looking for housing the next day are 20%. If she is partying, the chances of her partying the next day are 10%, the chances of her studying the next day are 60%, and the chances of her looking for housing the next day are 30%. If she is looking for housing, the chances of her looking for housing the next day are 50%, the chances of her partying the next day are 30% and the chances of her studying the next day are 20%.

(a) Draw a Markov chain to visualize Alex's life.

(b) Write out a matrix to represent this Markov chain.

(c) If Alex studies on Monday, what is the chance that she is partying on Friday?
(Don't do the math, just write out the expression that you would use to find it.)

(d) What percentage of her time should Alex expect to use looking for housing?

(e) If Alex parties on Monday, what is the chance of Alex partying again before studying?

2. Prehistoric States

A prehistoric civilization survives by hunting game in the forests near their home. At the beginning of the hunting season, all the young men go out to the forest. After the first day, those who have a kill, which happens with probability $1/2$, return home. Everyone who has been out for two days, even if without a kill, returns home for rest. And everyone who goes home goes back out the next day.

1. What are the states in this scenario? Draw a Markov chain.

2. What is the transition matrix? The initial vector?

3. Is this Markov chain reducible? Is it periodic?

4. What is the invariant vector?

5. What are the distributions after one week?

6. What is the expected length of hunting trip?

You have a database of an infinite number of movies. Each movie has a rating that is uniformly distributed in $0, 0.5, 1, 1.5, 2, 2.5, 3, 3.5, 4, 4.5, 5$ independent of all other movies. You want to find two movies such that the sum of their ratings is greater than 7.5 (7.5 is not included).

- Computer Science Mentors CS70 Spring 2017: Jerry Huang and Peijie Li, with David Harrison, Niraj Rao, Priya Agarwal, Katya Stukalova, Nikhil Dilip, Anwar Baroudi

4. **Bet On It**

Smith is in jail and has 3 dollars; he can get out on bail if he has 8 dollars. A guard agrees to make a series of bets with him. If Smith bets A dollars, he wins A dollars with probability 0.4 and loses A dollars with probability 0.6.

- a) Find the probability that he wins 8 dollars before losing all of his money if he bets 1 dollar each time.

- b) Find the probability that he wins 8 dollars before losing all of his money if he bets, each time, as much as possible but not more than necessary to bring his fortune up to 8 dollars

- c) Which strategy gives Smith the better chance of getting out of jail?

5. **Tossing Coins**

A fair coin is tossed repeatedly and independently. Find the expected number of tosses till the pattern HTH appears.